

Product Manager, Agentic Growth – Product Exercise

You'll own a compounding library of agentic systems — recurring, one-shot, and event-triggered — driving every Bold growth surface in parallel. **Your project goal: Build an agentic growth workflow; don't write a deck.**

Spend **≤4-5 hours**. Use whatever stack you use day-to-day; that fluency is part of what we're evaluating. We can cover costs of tools up to \$100 if needed. Please submit receipts to accounting@bold.org to get reimbursed.

Please send questions to tori@bold.org. Please send your work ≥24h before your interview to emily@mechanism.com, cc tori@bold.org, irina@bold.org, jason@bold.org, dror@bold.org. Please come to your interview ready to demo and discuss your project!

Context

Bold's top-of-funnel runs through three public surfaces — **scholarship detail pages**, **scholarship index / by-* pages (SGPs)**, and the **blog**.

Live pages to ground your thinking (target keyword italicized):

- [/scholarships/by-type/easy-scholarships](#) — *easy scholarships*
- [/scholarships/by-type/essay-scholarships](#) — *essay scholarships*
- [/scholarships/by-type/merit-based](#) — *merit based scholarships*
- [/scholarships/by-type/grants-to-pay-off-student-loans](#) — *grants to pay off student loans*
- [/scholarships/by-year/graduate-students-scholarships](#) — *graduate scholarships*
- [/scholarships/by-year/high-school/seniors](#) — *scholarships for high school seniors*
- [/scholarships/the-be-bold-no-essay-scholarship](#) — flagship detail page; very high CVR
- [/blog/best-physical-therapy-schools](#) — high-traffic blog that converts near zero

We encode growth work as **agentic systems**: recurring loops (title-tag optimization, paid bidding, lifecycle testing), one-shot generators, or event-triggered (e.g. new persona interviews land → SGP copy refresh + A/B test; one money page rebuilt end-to-end). We want to see your taste about which shape of a solution a growth problem deserves.

Your [reference dataset](#) includes funnel by surface, top landing pages with GSC data, device split, tech-SEO signals, instrumentation diagnostics, one live experiment. See the README sheet for details.

Part 1) Build (3–4h)

Pick the highest-leverage problem you can credibly solve with an agentic system in the time you have. Recurring loop, one-shot generator, or event-triggered — your call. **The choice itself is part of what we're evaluating**; defend it in your README.

Build it end-to-end: it should run, call a real LLM at least once, persist state where appropriate, and produce a useful artifact. Mocking gated APIs is fine.

Deliverables: runnable system (link); workflow artifact (skill/playbook/prompt); sample output; ≤1-page README covering what you chose and why, what you'd do with another day / week, the riskiest part, and how you'd earn trust before it acts without human review.

Part 2) Design two more (1–2h)

Sketch two systems **complementary to Part 1**. At least one should be a different shape (if Part 1 is a loop, this one isn't).

≤1 page each: what it does, trigger, data/tools, state between runs (if any), decision and guardrail rules, how you'd know it's working, how you'd graduate it toward acting without you.

Part 3) Avoiding fake wins (≤30m)

Short prose: how would you detect instrumentation bugs, traffic-mix shifts, and downstream quality regressions (fake/low-intent signups, low Day-7 activation)? What guardrails would you hard-code into the systems?